

# DATA 621—Assignment no. 1

*Critical Thinking Group 2: All of our names, Ben H.*

*September XX, 2019*

## Contents

<b>Executive Overview</b>	<b>1</b>
<b>Data Exploration</b>	<b>2</b>
Response Variable: Yearly wins . . . . .	2
Explanatory Variables . . . . .	4
<b>Data Preparation</b>	<b>9</b>
<b>Modeling</b>	<b>10</b>
Model 1 . . . . .	10
Model 2 . . . . .	11
Model 3 . . . . .	12
<b>Evaluation</b>	<b>13</b>
Model 1 . . . . .	13
Model 2 . . . . .	16
Model 3 . . . . .	17
<b>Predictions</b>	<b>18</b>

Introduction

Load libraries:

```
library(corrplot)
library(dplyr)
library(ggplot2)
library(knitr)

library(kableExtra)
# library(psych)
library(reshape2)
library(gridExtra)
#library(psychometric)
#library(ggpubr)
#library(matlib)
#library(matrixcalc)
# library(psych)
# library(MASS)
# library(forecast)
```

## Executive Overview

We present three multiple regression models to predict a professional baseball teams' performance.

# Data Exploration

The training data has 17 columns and 2,276 rows.

The explanatory columns are broken down into four categories:

- Batting
- Base run
- Pitching
- Fielding

Below you will see a preview of the columns and the first few observations broken down into these four categories.

##	INDEX	TARGET_WINS	TEAM_BATTING_H	TEAM_BATTING_2B	TEAM_BATTING_3B
## 1	1	39	1445	194	39
## 2	2	70	1339	219	22
## 3	3	86	1377	232	35
## 4	4	70	1387	209	38
## 5	5	82	1297	186	27
## 6	6	75	1279	200	36
##	TEAM_BATTING_HR	TEAM_BATTING_BB	TEAM_BATTING_SO	TEAM_BASERUN_SB	
## 1	13	143	842	NA	
## 2	190	685	1075	37	
## 3	137	602	917	46	
## 4	96	451	922	43	
## 5	102	472	920	49	
## 6	92	443	973	107	
##	TEAM_BASERUN_CS	TEAM_BATTING_HBP	TEAM_PITCHING_H	TEAM_PITCHING_HR	
## 1	NA	NA	9364	84	
## 2	28	NA	1347	191	
## 3	27	NA	1377	137	
## 4	30	NA	1396	97	
## 5	39	NA	1297	102	
## 6	59	NA	1279	92	
##	TEAM_PITCHING_BB	TEAM_PITCHING_SO	TEAM_FIELDING_E	TEAM_FIELDING_DP	
## 1	927	5456	1011	NA	
## 2	689	1082	193	155	
## 3	602	917	175	153	
## 4	454	928	164	156	
## 5	472	920	138	168	
## 6	443	973	123	149	

## Response Variable: Yearly wins


The variable **TARGET\_WINS** is the number of wins of a professional baseball team for a given year. The year is not part of the data set.


This is the dependent variable that our models will attempt to predict. It is characterized by:

TARGET_WINS
Min. : 0.00
1st Qu.: 71.00
Median : 82.00
Mean : 80.79
3rd Qu.: 92.00
Max. :146.00

	n	sd	se
TARGET_WINS	2276	15.75215	0.3301823

the distribution of the number of wins is unimodal and skewed to the left with some outliers towards the tail. It looks approximately normal, though the boxplot shows there are quite a few outliers. The minimum number of wins for a team is 0 and the maximum is 146. The mean is 80.79.






## Explanatory Variables

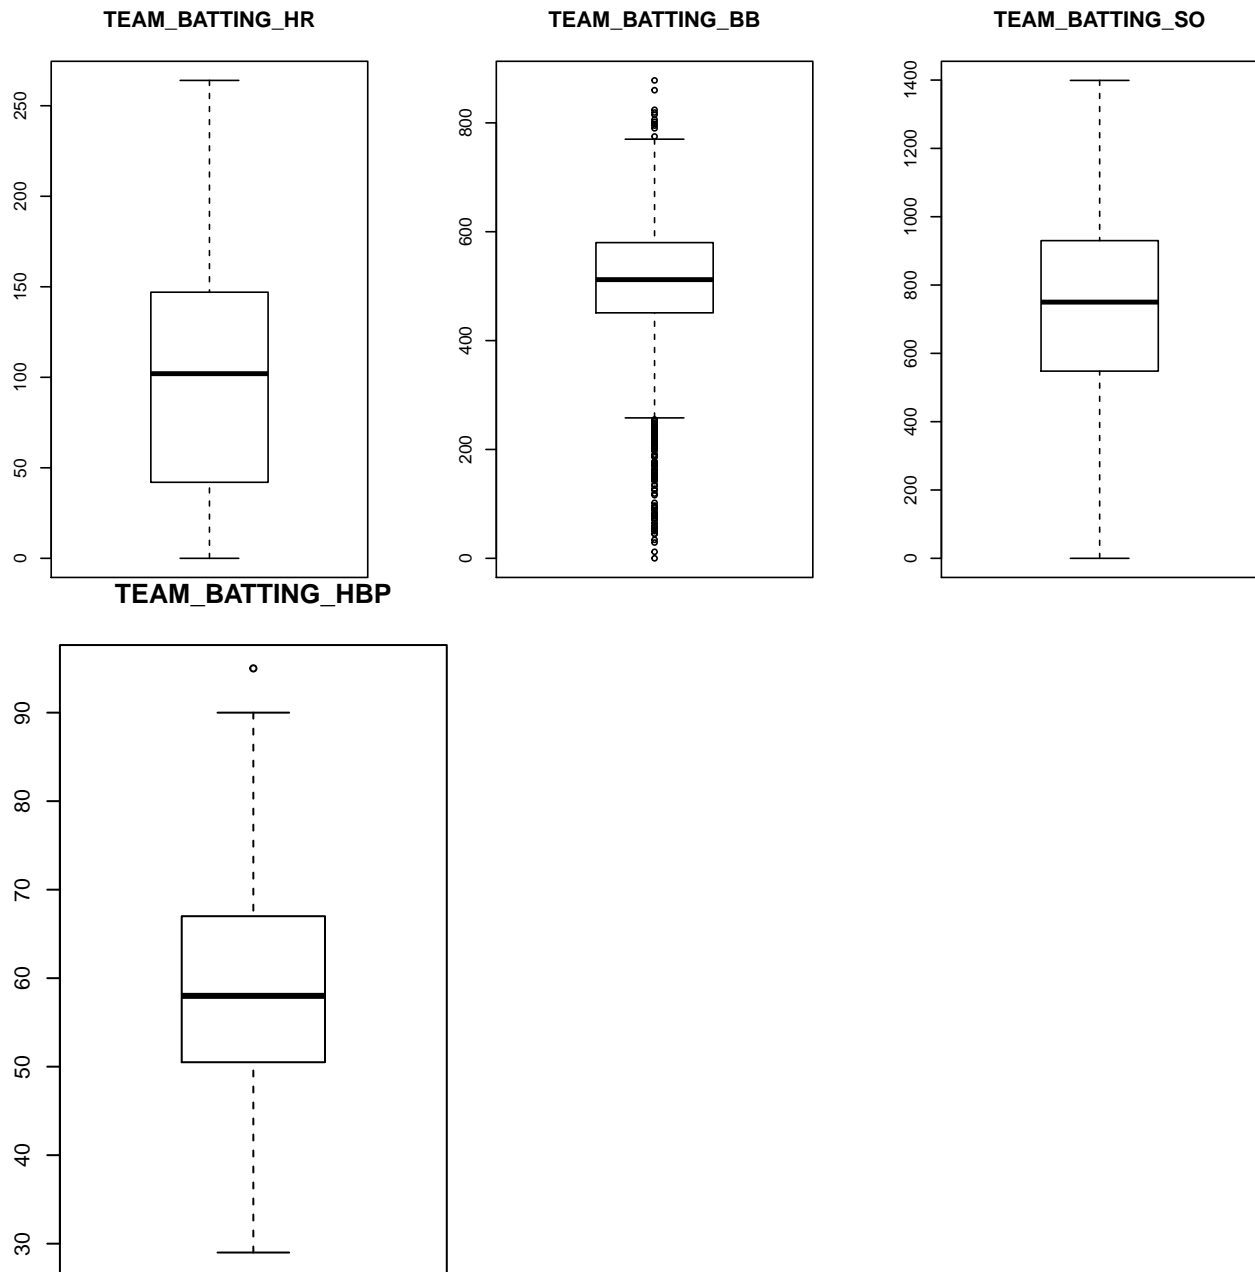
### Batting variables

Below are our batting variables and their hypothesized effect on **TARGET\_WINS**:

DESCRIPTION AND THEORETICAL EFFECT

VARIABLE NAME	DEFINITION	THEORETICAL EFFECT
TEAM_BATTING_H	HBase Hits by batters (1B,2B,3B,HR)	Positive Impact on Wins
TEAM_BATTING_2B	Doubles by batters (2B)	Positive Impact on Wins
TEAM_BATTING_3B	Triples by batters (3B)	Positive Impact on Wins
TEAM_BATTING_HR	Homeruns by batters (4B)	Positive Impact on Wins
TEAM_BATTING_BB	Walks by batters	Positive Impact on Wins
TEAM_BATTING_HBP	Batters hit by pitch (get a free base)	Positive Impact on Wins
TEAM_BATTING_SO	Strikeouts by batters	Negative Impact on Wins





The boxplots hint that some of these variables are quite skewed, especially `TEAM_BATTING_BB` and `TEAM_BATTING_HBP`.

Since all of these variables relate to the same thing, batting, we expect at least some of them to be correlated. This has implications on later modeling:

	TARGET_WINS	TEAM_BATTING_HR	TEAM_BATTING_BB	TEAM_BATTING_SO	TEAM_BATTING_HBP
TARGET_WINS	1.000000	0.4699467	0.3129840	-	0.0735042
TEAM_BATTING_HR	0.4699467	1.0000000	0.1243459	0.2288927	-
TEAM_BATTING_BB	0.3129840	0.1243459	1.0000000	0.3962759	0.0291122
TEAM_BATTING_SO	0.2288927	0.3962759	0.3962759	1.0000000	0.0460848
TEAM_BATTING_HBP	0.0735042	-	0.0291122	0.0460848	1.0000000

[illegible]

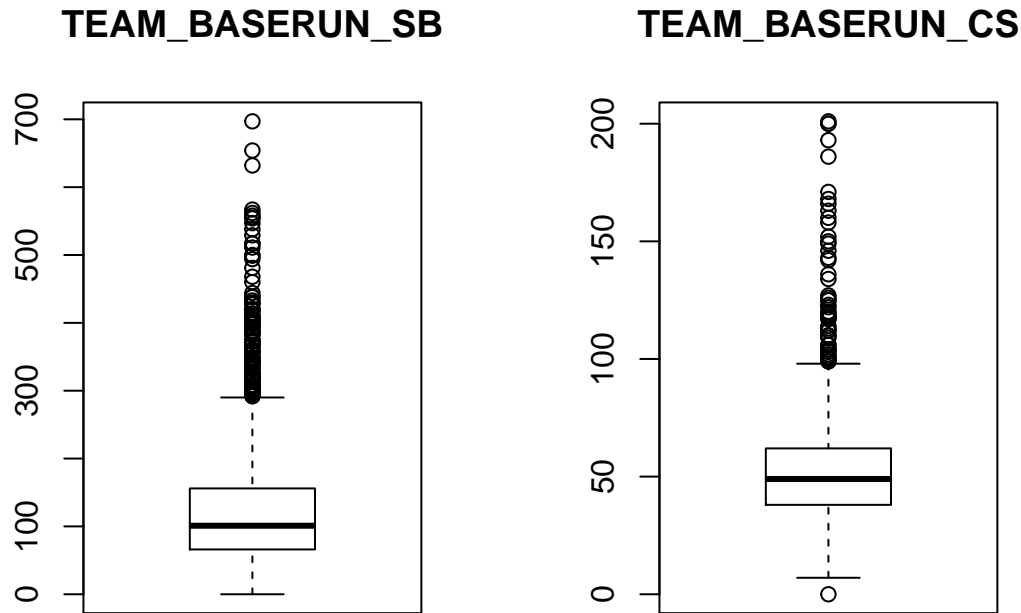
## Baserun Variables

Description and theoretical effects:

VARIABLE NAME	DEFINITION	THEORETICAL EFFECT
TEAM_BASERUN_SB	Stolen bases	Positive Impact on Wins
TEAM_BASERUN_CS	Caught stealing	Negative Impact on Wins

As you can see, both variables have some missing values:

TEAM_BASERUN_SB	TEAM_BASERUN_CS
Min. : 0.0	Min. : 0.0
1st Qu.: 66.0	1st Qu.: 38.0
Median :101.0	Median : 49.0
Mean :124.8	Mean : 52.8
3rd Qu.:156.0	3rd Qu.: 62.0
Max. :697.0	Max. :201.0
NA's :131	NA's :772



	TARGET_WINS	TEAM_BASERUN_SB	TEAM_BASERUN_CS
TARGET_WINS	1.0000000	0.1539213	0.0224041
TEAM_BASERUN_SB	0.1539213	1.0000000	0.6552448
TEAM_BASERUN_CS	0.0224041	0.6552448	1.0000000

## Pitching Variables

Description:

VARIABLE NAME	DEFINITION	THEORETICAL EFFECT
TEAM_PITCHING_BB	Walks allowed	Negative Impact on Wins
TEAM_PITCHING_H	Hits allowed	Negative Impact on Wins
TEAM_PITCHING_HR	Homeruns allowed	Negative Impact on Wins
TEAM_PITCHING_SO	Strikeouts by pitchers	Positive Impact on Wins

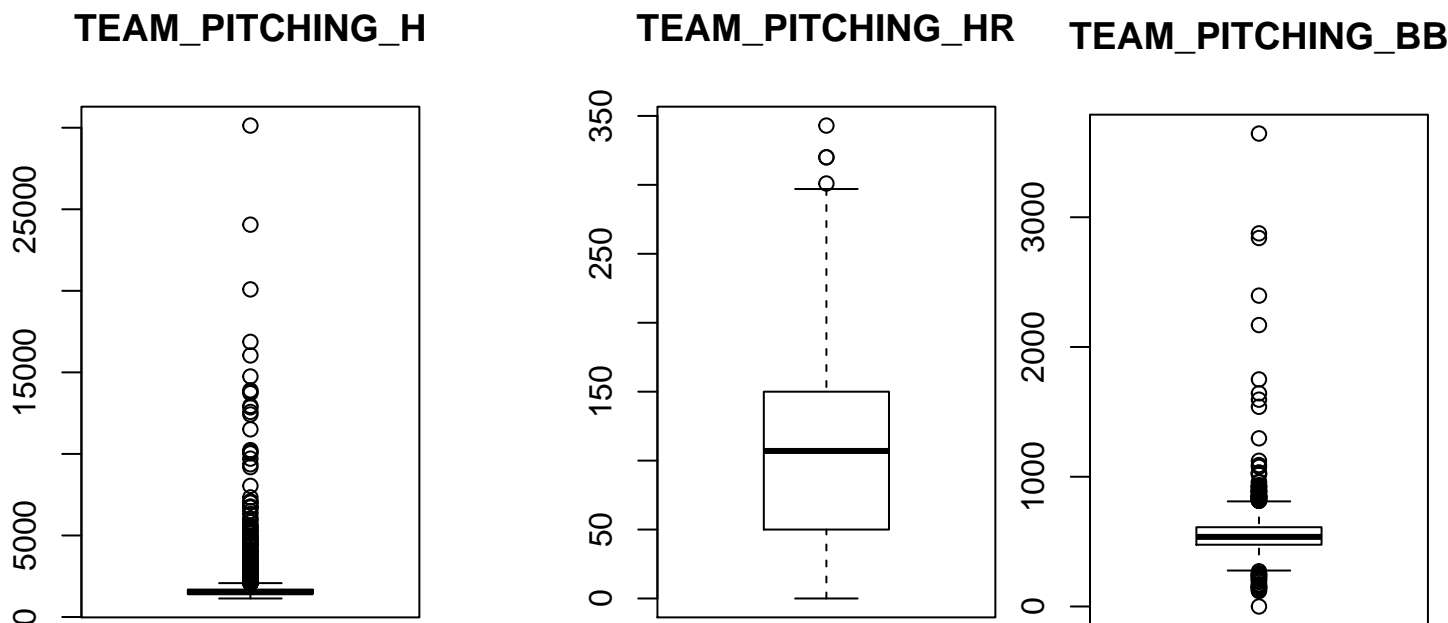
```
kable(head(data[c(12:15)]), format='markdown')
```

TEAM_PITCHING_H	TEAM_PITCHING_HR	TEAM_PITCHING_BB	TEAM_PITCHING_SO
9364	84	927	5456
1347	191	689	1082
1377	137	602	917
1396	97	454	928
1297	102	472	920
1279	92	443	973

```
kable(summary(data[c(12:15)]), format='markdown')
```

TEAM_PITCHING_H	TEAM_PITCHING_HR	TEAM_PITCHING_BB	TEAM_PITCHING_SO
Min. : 1137	Min. : 0.0	Min. : 0.0	Min. : 0.0

TEAM_PITCHING_H	TEAM_PITCHING_HR	TEAM_PITCHING_BB	TEAM_PITCHING_SO
1st Qu.: 1419	1st Qu.: 50.0	1st Qu.: 476.0	1st Qu.: 615.0
Median : 1518	Median :107.0	Median : 536.5	Median : 813.5
Mean : 1779	Mean :105.7	Mean : 553.0	Mean : 817.7
3rd Qu.: 1682	3rd Qu.:150.0	3rd Qu.: 611.0	3rd Qu.: 968.0
Max. :30132	Max. :343.0	Max. :3645.0	Max. :19278.0
NA	NA	NA	NA's :102



```
kable(cor(data[c(2, 12:15)]), format='markdown')
```

	TARGET_WINS	TEAM_PITCHING_H	TEAM_PITCHING_HR	TEAM_PITCHING_BB	TEAM_PITCHING_SO
TARGET_WINS	1.0000000	-0.1099371	0.1890137	0.1241745	NA
TEAM_PITCHING_H	-0.1099371	1.0000000	-0.1416128	0.3206762	NA
TEAM_PITCHING_HR	-0.1416128	-0.1416128	1.0000000	0.2219375	NA
TEAM_PITCHING_BB	0.1241745	0.3206762	0.2219375	1.0000000	NA
TEAM_PITCHING_SO	NA	NA	NA	NA	1

## Fielding Variables

Description:

VARIABLE NAME	DEFINITION	THEORETICAL EFFECT
TEAM_FIELDING_E	Errors	Negative Impact on Wins
TEAM_FIELDING_DP	Double Plays	Positive Impact on Wins


```
kable(summary(data[c(16:17)]), format='markdown')
```

TEAM_FIELDING_E	TEAM_FIELDING_DP
Min. : 65.0	Min. : 52.0




TEAM_FIELDING_E	TEAM_FIELDING_DP
1st Qu.: 127.0	1st Qu.:131.0
Median : 159.0	Median :149.0
Mean : 246.5	Mean :146.4
3rd Qu.: 249.2	3rd Qu.:164.0
Max. :1898.0	Max. :228.0
NA	NA's :286

## TEAM\_FIELDING\_E



## TEAM\_FIELDING\_DP



```
kable(cor(data[c(2, 16:17)]))
```

	TARGET_WINS	TEAM_FIELDING_E	TEAM_FIELDING_DP
TARGET_WINS	1.0000000	-0.1764848	NA
TEAM_FIELDING_E	-0.1764848	1.0000000	NA
TEAM_FIELDING_DP	NA	NA	1

## Data Preparation

There are missing data in this dataset:

```
sort(sapply(data, FUN=function(x) mean(is.na(x))), decreasing=TRUE)
```

```
## TEAM_BATTING_HBP TEAM_BASERUN_CS TEAM_FIELDING_DP TEAM_BASERUN_SB
##      0.91608084      0.33919156      0.12565905      0.05755712
## TEAM_BATTING_SO TEAM_PITCHING_SO      INDEX      TARGET_WINS
##      0.04481547      0.04481547      0.00000000      0.00000000
## TEAM_BATTING_H TEAM_BATTING_2B TEAM_BATTING_3B TEAM_BATTING_HR
##      0.00000000      0.00000000      0.00000000      0.00000000
## TEAM_BATTING_BB TEAM_PITCHING_H TEAM_PITCHING_HR TEAM_PITCHING_BB
##      0.00000000      0.00000000      0.00000000      0.00000000
## TEAM_FIELDING_E
##      0.00000000
```

The TEAM\_BATTING\_HBP variable is almost entirely missing. We chose to omit it from the rest of the analysis.

As for the other variables, we are simply dropping them, further analyzing only the complete cases.

```
data$TEAM_BATTING_HBP <- NULL
data2 <- data[complete.cases(data), ]
```

Some variables can be transformed to approximately normal distribution by logging them. We will test this below:

```
train <- data
train.df <- data
train.trans <- train.df
# cols.nomiss = names(train.df)[!names(train.df)%in%(cols.miss)]
train.trans$TEAM_PITCHING_HR = apply(train.trans['TEAM_PITCHING_HR'], 1, function(x) if(x==0){return(0)}
train.trans$TEAM_PITCHING_BB = apply(train.trans['TEAM_PITCHING_BB'], 1, function(x) if(x==0){return(0)}
train.trans$TEAM_FIELDING_E = apply(train.trans['TEAM_FIELDING_E'], 1, function(x) if(x==0){return(0)} <
```

## Modeling

```
# SAMRITA's prep ?
train.trans$TEAM_BATTING_H = train.trans$TEAM_BATTING_H + train.trans$TEAM_BATTING_2B + train.trans$TEAM_BATTING_3B
train.trans$TEAM_FIELDING_E = train.trans$TEAM_FIELDING_E + train.trans$TEAM_PITCHING_H
train.trans = train.trans%>%dplyr::select(-TEAM_BATTING_2B, -TEAM_BATTING_3B, -TEAM_BATTING_HR, -TEAM_PITCHING_HR)
```

### Model 1

The first model regresses `target_wins` on the five variables identified above:

```
model1 = lm(TARGET_WINS ~ TEAM_BATTING_H + TEAM_BATTING_BB + TEAM_PITCHING_HR +
            TEAM_PITCHING_BB + TEAM_FIELDING_E, data=train.trans)
summary(model1)
```

```
##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BATTING_H + TEAM_BATTING_BB +
##     TEAM_PITCHING_HR + TEAM_PITCHING_BB + TEAM_FIELDING_E, data = train.trans)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -50.087  -9.176   0.455   9.110  59.641
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   2.8223840   8.3128468   0.340    0.734
## TEAM_BATTING_H  0.0383705   0.0017159  22.362 < 2e-16 ***
## TEAM_BATTING_BB  0.0167647   0.0041036   4.085 4.55e-05 ***
## TEAM_PITCHING_HR -1.9009586   0.4477722  -4.245 2.27e-05 ***
## TEAM_PITCHING_BB  1.5483525   1.5160674   1.021   0.307
## TEAM_FIELDING_E -0.0018848   0.0002602  -7.244 5.95e-13 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.77 on 2270 degrees of freedom
## Multiple R-squared:  0.2377, Adjusted R-squared:  0.236
```

```
## F-statistic: 141.6 on 5 and 2270 DF,  p-value: < 2.2e-16
```

All but one of the variables are extremely significant, including the intercept parameter. The  $F$ -statistic suggests the model is not just picking up random noise.

It estimates that base hits and walks by batters are positively related to team wins, and that walks allowed and errors and homeruns allowed are negatively related. Curiously, it finds `TEAM_PITCHING_BB` is positively related, although it was theorized to be negatively related.

Mean squared error, as a baseline for evaluation:

```
mean(resid(model1)^2)
```

```
## [1] 189.0616
```

## Model 2

The second model uses the same variables, but some of them have been log-transformed, and we use squared functions of some of them, rather than the variable itself.

```
train.trans = train.df
train.trans$TEAM_BATTING_H = train.trans$TEAM_BATTING_H + train.trans$TEAM_BATTING_2B + train.trans$TEAM_BATTING_BB
train.trans$TEAM_FIELDING_E = train.trans$TEAM_FIELDING_E + train.trans$TEAM_PITCHING_H
train.trans = train.trans%>%dplyr::select(-TEAM_BATTING_2B, -TEAM_BATTING_3B, -TEAM_BATTING_HR, -TEAM_PITCHING_BB)
summary(train.trans)
```

```
##      INDEX      TARGET_WINS      TEAM_BATTING_H TEAM_BATTING_BB
##  Min.   : 1.0    Min.   : 0.00    Min.   :1026    Min.   : 0.0
## 1st Qu.: 630.8  1st Qu.: 71.00    1st Qu.:1739    1st Qu.:451.0
## Median :1270.5  Median : 82.00    Median :1862    Median :512.0
## Mean   :1268.5  Mean   : 80.79    Mean   :1865    Mean   :501.6
## 3rd Qu.:1915.5  3rd Qu.: 92.00    3rd Qu.:1978    3rd Qu.:580.0
## Max.   :2535.0  Max.   :146.00    Max.   :3092    Max.   :878.0
##
## TEAM_BATTING_SO TEAM_BASERUN_SB TEAM_BASERUN_CS TEAM_PITCHING_HR
##  Min.   : 0.0    Min.   : 0.0    Min.   : 0.0    Min.   : 0.0
## 1st Qu.: 548.0  1st Qu.: 66.0    1st Qu.: 38.0    1st Qu.: 50.0
## Median : 750.0  Median :101.0    Median : 49.0    Median :107.0
## Mean   : 735.6  Mean   :124.8    Mean   : 52.8    Mean   :105.7
## 3rd Qu.: 930.0  3rd Qu.:156.0    3rd Qu.: 62.0    3rd Qu.:150.0
## Max.   :1399.0  Max.   :697.0    Max.   :201.0    Max.   :343.0
## NA's   :102    NA's   :131     NA's   :772
## TEAM_PITCHING_BB TEAM_PITCHING_SO TEAM_FIELDING_E TEAM_FIELDING_DP
##  Min.   : 0.0    Min.   : 0.0    Min.   :1276    Min.   : 52.0
## 1st Qu.: 476.0  1st Qu.: 615.0    1st Qu.:1566    1st Qu.:131.0
## Median : 536.5  Median : 813.5    Median :1679    Median :149.0
## Mean   : 553.0  Mean   : 817.7    Mean   :2026    Mean   :146.4
## 3rd Qu.: 611.0  3rd Qu.: 968.0    3rd Qu.:1922    3rd Qu.:164.0
## Max.   :3645.0  Max.   :19278.0    Max.   :31860    Max.   :228.0
## NA's   :102    NA's   :102     NA's   :286
```

```
train.trans$TEAM_PITCHING_HR = apply(train.trans[, 'TEAM_PITCHING_HR'], 1, function(x) if(x==0){return(0)} else {sqrt(x)})
train.trans$TEAM_PITCHING_BB = apply(train.trans[, 'TEAM_PITCHING_BB'], 1, function(x) if(x==0){return(0)} else {sqrt(x)})
train.trans$TEAM_FIELDING_E = apply(train.trans[, 'TEAM_FIELDING_E'], 1, function(x) if(x==0){return(0)} else {sqrt(x)})
```

```
model2 = lm(TARGET_WINS ~ TEAM_BATTING_H + poly(TEAM_BATTING_BB,2) +
            TEAM_PITCHING_HR + poly(TEAM_PITCHING_BB,2) +
```

```

poly(Team_FIELDING_E,2), data=train.trans)
summary(model2)

##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BATTING_H + poly(Team_BATTING_BB,
##      2) + TEAM_PITCHING_HR + poly(Team_PITCHING_BB, 2) + poly(Team_FIELDING_E,
##      2), data = train.trans)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -52.787  -8.938   0.192   9.165  49.336
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.848e+00  3.764e+00   0.757   0.449
## TEAM_BATTING_H      4.945e-02  2.586e-03  19.122 < 2e-16 ***
## poly(Team_BATTING_BB, 2)1 -3.781e+02  6.711e+01  -5.634 1.98e-08 ***
## poly(Team_BATTING_BB, 2)2   2.057e+02  2.629e+01   7.826 7.65e-15 ***
## TEAM_PITCHING_HR    -3.246e+00  5.333e-01  -6.088 1.34e-09 ***
## poly(Team_PITCHING_BB, 2)1  2.847e+02  4.062e+01   7.009 3.16e-12 ***
## poly(Team_PITCHING_BB, 2)2  1.759e+02  2.884e+01   6.100 1.24e-09 ***
## poly(Team_FIELDING_E, 2)1  -5.730e+02  6.423e+01  -8.920 < 2e-16 ***
## poly(Team_FIELDING_E, 2)2  -1.043e+02  1.606e+01  -6.496 1.01e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.58 on 2267 degrees of freedom
## Multiple R-squared:  0.259, Adjusted R-squared:  0.2563
## F-statistic: 99.02 on 8 and 2267 DF,  p-value: < 2.2e-16

```

All of our parameters are significant at  $p < 0.001$  level, although the intercept term is no longer as significant as it was. Adjusted  $R^2$  has improved, as has  $MSE$ :

```
mean(resid(model1))^2
```

```
## [1] 189.0616
```

## Model 3

model3 attempts to more precisely model TEAM\_FIELDING\_E by cubing it.

```

model3 = lm(TARGET_WINS ~ TEAM_BATTING_H + poly(Team_BATTING_BB,2) +
            TEAM_PITCHING_HR + poly(Team_PITCHING_BB,2) +
            poly(Team_FIELDING_E,3), data=train.trans)
summary(model3)

##
## Call:
## lm(formula = TARGET_WINS ~ TEAM_BATTING_H + poly(Team_BATTING_BB,
##      2) + TEAM_PITCHING_HR + poly(Team_PITCHING_BB, 2) + poly(Team_FIELDING_E,
##      3), data = train.trans)
##
## Residuals:
##      Min       1Q   Median       3Q      Max

```

```
## -51.800 -8.957 0.291 9.101 49.593
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      2.836e+00  3.756e+00  0.755  0.45028
## TEAM_BATTING_H      5.083e-02  2.614e-03 19.443 < 2e-16 ***
## poly(Team_BATTING_BB, 2)1 -3.519e+02  6.744e+01 -5.218 1.97e-07 ***
## poly(Team_BATTING_BB, 2)2  1.826e+02  2.716e+01  6.722 2.26e-11 ***
## TEAM_PITCHING_HR    -3.827e+00  5.607e-01 -6.826 1.12e-11 ***
## poly(Team_PITCHING_BB, 2)1  2.781e+02  4.058e+01  6.855 9.19e-12 ***
## poly(Team_PITCHING_BB, 2)2  1.952e+02  2.937e+01  6.646 3.77e-11 ***
## poly(Team_FIELDING_E, 3)1 -5.658e+02  6.413e+01 -8.823 < 2e-16 ***
## poly(Team_FIELDING_E, 3)2 -1.051e+02  1.602e+01 -6.561 6.61e-11 ***
## poly(Team_FIELDING_E, 3)3 -5.528e+01  1.682e+01 -3.286 0.00103 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.55 on 2266 degrees of freedom
## Multiple R-squared:  0.2625, Adjusted R-squared:  0.2595
## F-statistic: 89.6 on 9 and 2266 DF, p-value: < 2.2e-16
```

All the model parameters are highly significant, although the intercept is no longer significant. This is fine, however; arguably, a team with scores of 0 across all variables could expect to have zero wins on average.

Adjusted  $R^2$  and  $MSE$  tick up slightly compared to previous models:

```
mean(resid(model13)^2)
```

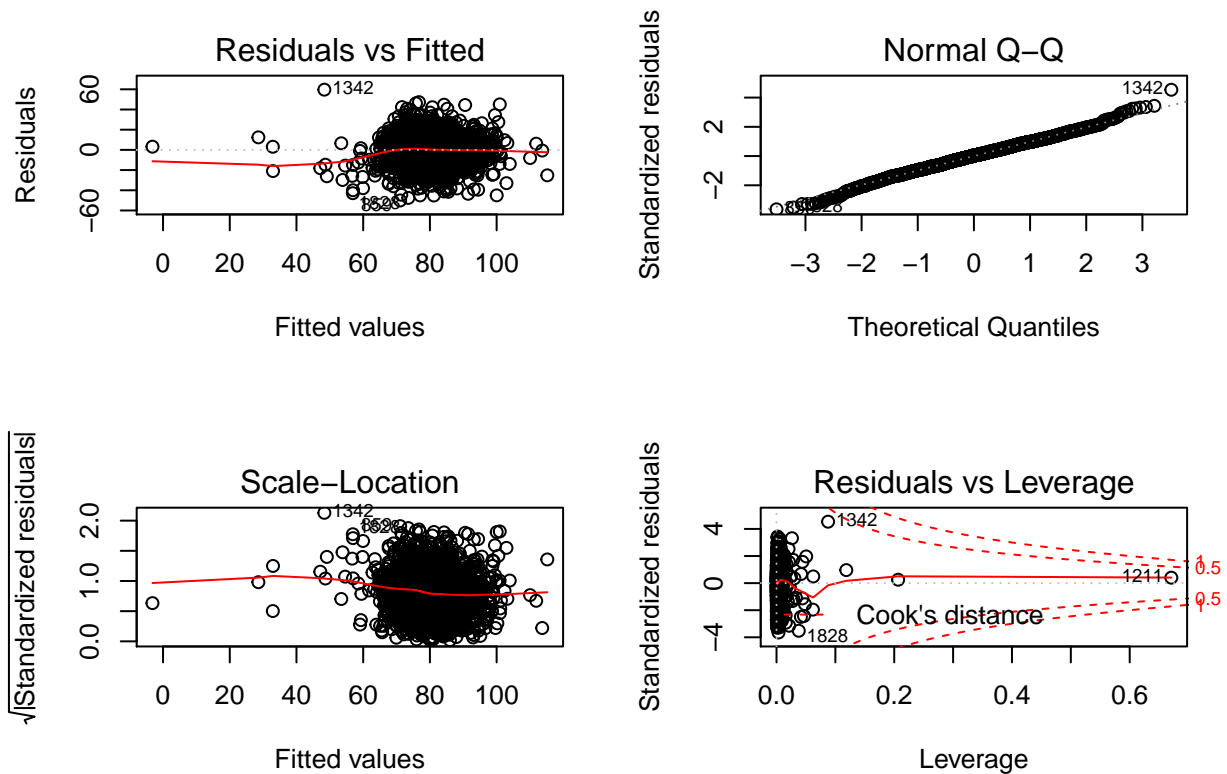
```
## [1] 182.9232
```

## Evaluation

This section evaluates the above models, particularly focusing on residuals analysis.

### Model 1

```
par(mfrow=c(2,2))
plot(model1)
```



Although the Q-Q and histogram plots show the residuals aren't too bad, there does appear to be some systematic bias. Lower fitted values from approximately 20-60 are systematically negative.

There's a strange point in the residuals plot, around  $\hat{y} = 15$ . Examining this point more closely, we don't see any obvious cause for it:

```
train[which.min(resid(model1)), ]
```

```
##      INDEX TARGET_WINS TEAM_BATTING_H TEAM_BATTING_2B TEAM_BATTING_3B
## 859    950          21         1402         149         53
##      TEAM_BATTING_HR TEAM_BATTING_BB TEAM_BATTING_SO TEAM_BASERUN_SB
## 859           13         304         295         134
##      TEAM_BASERUN_CS TEAM_PITCHING_H TEAM_PITCHING_HR TEAM_PITCHING_BB
## 859           NA        1475         14         320
##      TEAM_PITCHING_SO TEAM_FIELDING_E TEAM_FIELDING_DP
## 859          310         408           NA
```

This point has extremely high leverage:

```
train[1211, ]
```

```
##      INDEX TARGET_WINS TEAM_BATTING_H TEAM_BATTING_2B TEAM_BATTING_3B
## 1211   1347          0         891         135          0
##      TEAM_BATTING_HR TEAM_BATTING_BB TEAM_BATTING_SO TEAM_BASERUN_SB
## 1211          0          0          0          0
##      TEAM_BASERUN_CS TEAM_PITCHING_H TEAM_PITCHING_HR TEAM_PITCHING_BB
## 1211          0        24057          0          0
##      TEAM_PITCHING_SO TEAM_FIELDING_E TEAM_FIELDING_DP
## 1211          0        1890           NA
```




It is the only team with 0 wins in the entire dataset, and naturally the model could not accurately estimate wins for this data point. Future modeling should consider excluding this unusual case.

Look at each variable plotted against the model's residuals to see if we can understand the source of some of this bias:

```
par(mfrow=c(3,1))
plot(train$TEAM_BATTING_H, resid(model1))
abline(lm( resid(model1) ~ train$TEAM_BATTING_H))



plot(train$TEAM_BATTING_BB, resid(model1))
abline(lm( resid(model1) ~ train$TEAM_BATTING_BB))

plot(train$TEAM_PITCHING_HR, resid(model1))
abline(lm( resid(model1) ~ train$TEAM_PITCHING_HR))
```



```
par(mfrow=c(2,1))
plot(train$TEAM_PITCHING_BB, resid(model1))
abline(lm( resid(model1) ~ train$TEAM_PITCHING_BB))

plot(train$TEAM_FIELDING_E, resid(model1))
abline(lm( resid(model1) ~ train$TEAM_FIELDING_E))
```



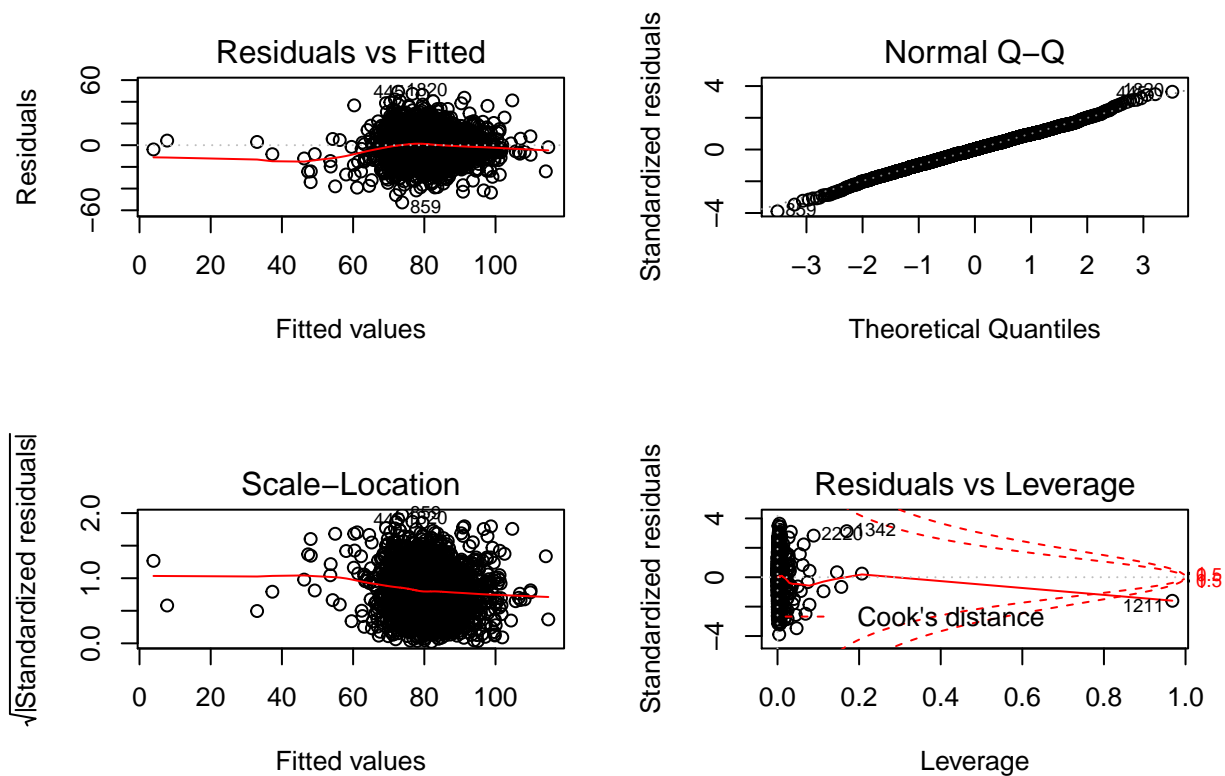
The trend lines all look fairly level, the least squares' assumption of constant variance is violated in most of these plots. Additionally, the model consistently underestimates when `TEAM_FIELDING_E` is around or greater than 1500.

## Model 2

From above, we know that this model performs better than `model11`. Examine its residuals:

```
par(mfrow=c(2,2))
plot(model12)
```





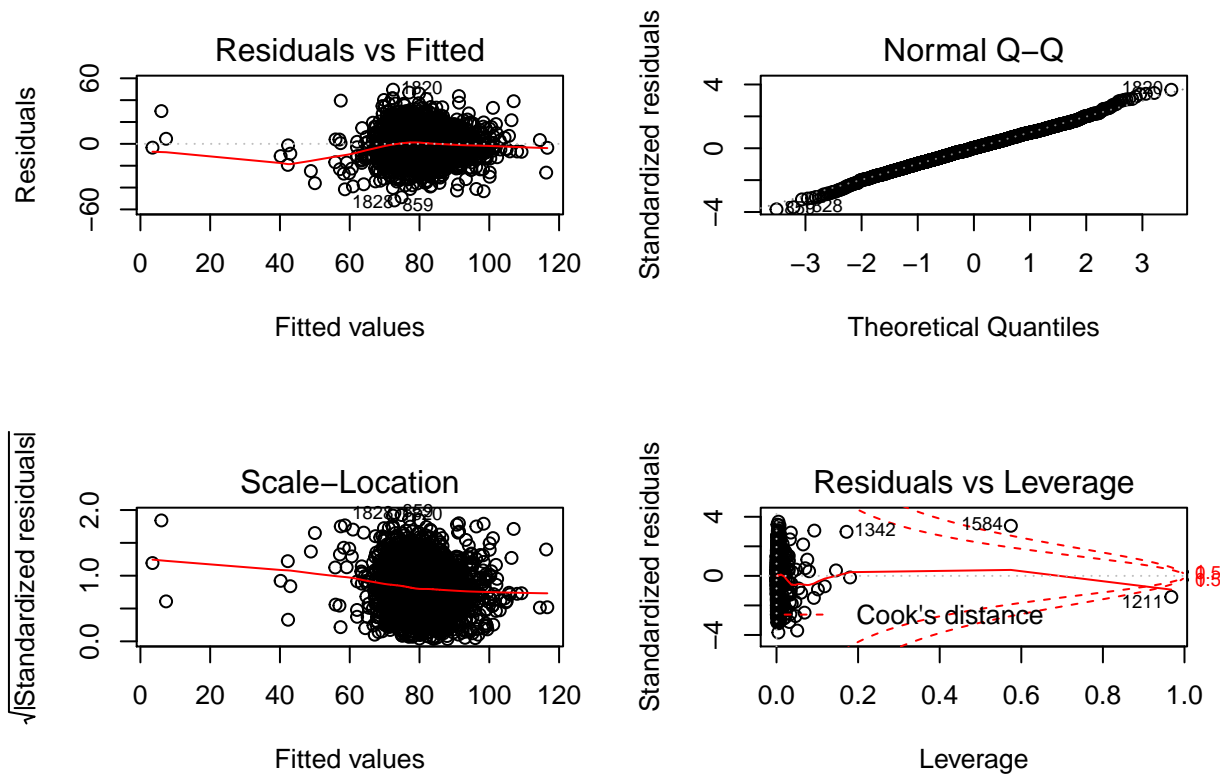
Thanks to the transformations, the Q-Q plot indicates these residuals are somewhat more normal than before. However, despite the transformations, data point 1211 is still an issue.

Systematic bias in the model remains, although it appears to be on the other side: as the predicted `TARGET_WINS` increases, the residuals are more strongly negative. I.e., this model is under-predicting better performing teams.

### Model 3

Finally, model three:

```
par(mfrow=c(2,2))
plot(model3)
```



Although this model had the best adjusted  $R^2$  and  $MSE$ , it appears even more biased than the previous model, but in the same way.

The plots of each variable against the model residual suggests the same problems plague model3 as the others: Unusual outlying points and lack of constant variance.

## Predictions

```
df.test = read.table("moneyball-evaluation-data.csv", sep=',', header = TRUE, stringsAsFactors = FALSE)
#2] SELECT Non Null columns
df.test = df.test %>% dplyr::select(-INDEX)
#cols.nomiss = names(train.df)[!names(train.df)%in%(cols.miss)]
#df.test = df.test %>% dplyr::select(cols.nomiss)
#3] Transform Data
test.trans = df.test
test.trans$TEAM_BATTING_H = test.trans$TEAM_BATTING_H + test.trans$TEAM_BATTING_2B + test.trans$TEAM_BATTING_3B
test.trans$TEAM_FIELDING_E = test.trans$TEAM_FIELDING_E + test.trans$TEAM_PITCHING_H
test.trans = test.trans %>% dplyr::select(-TEAM_BATTING_2B, -TEAM_BATTING_3B, -TEAM_BATTING_HR, -TEAM_PITCHING_HR)
test.trans$TEAM_PITCHING_HR = apply(test.trans['TEAM_PITCHING_HR'], 1, function(x) if(x==0){return(0)} else{return(x)})
test.trans$TEAM_PITCHING_BB = apply(test.trans['TEAM_PITCHING_BB'], 1, function(x) if(x==0){return(0)} else{return(x)})
test.trans$TEAM_FIELDING_E = apply(test.trans['TEAM_FIELDING_E'], 1, function(x) if(x==0){return(0)} else{return(x)})
Target_Wins = predict(model3, newdata = test.trans)
test.trans$TARGET_WINS = Target_Wins
df.test1 = read.table("moneyball-evaluation-data.csv", sep=',', header = TRUE, stringsAsFactors = FALSE)
df.test1$TARGET_WINS = Target_Wins
write.csv(df.test1, "Target_Prediction.csv", row.names = FALSE)
```